

Data Science Capstone Milestone

R. Martinez

February 19, 2019

Introduction

This is Capstone project for the Data Science Specialization of the John Hopkins University. The main objective is to build a text model to predict the next words that a user will type on a text device such a mobile text message app.

This document shows the milestone report for week #2 where I will:

- Get the data file to train the text model.
- Perform data exploration to understand the data.
- Profile the data file to identify basic characteristics of the data.
- Provide a initial plan to build the prediction text model.

.

For detail about the code, here is the link to the GitHub repo: <https://github.com/grauml/Data-Science-Capstone>.

.

Library and Data

The data was provided by the professor and can be downloaded from this source: <https://d396qusza40orc.cloudfront.net/dsscystone/dataset/Coursera-SwiftKey.zip>

```
library(tm); library(SnowballC); library(wordcloud); library(dplyr)
```

Loading required package: NLP

Loading required package: RColorBrewer

Attaching package: 'dplyr'

The following objects are masked from 'package:stats':

```
filter, lag
```

The following objects are masked from 'package:base':

```
intersect, setdiff, setequal, union
```

```
library(RWeka); library(ggplot2); library(formattable); library(tokenizers)
```

Attaching package: 'ggplot2'

The following object is masked from 'package:NLP':

```
annotate
```

```
set.seed(201902)
```

```
# Note that due the size of the data files (200MG, 196MG, and 159MG)
```

```
# I will use smaller size of the files selecting lines randomly (using rbinom)
```

```
#file_path <- "C:/Users/raulm/Documents/John Hopkins University/#10 Data Science Capstone/Coursera-Swi.
```

```

file_path <- "C:/Users/rlmartinez/Documents/John Hopkins University/#10 Data Science Capstone/Coursera
file_blogs <- "en_US.blogs.txt"
file_news <- "en_US.news.txt"
file_twitter <- "en_US.twitter.txt"

# Load the "blogs" data file
incon <- file(paste(file_path, file_blogs ,sep=""), "r")
file <- readLines(incon, encoding="UTF-8", skipNul=TRUE)
data_blogs <- file[rbinom(length(file), 1, 0.01) == 1]
close(incon)
len_blogs <- length(file)
size_blogs <- object.size(file)
max_line_blogs <- max(nchar(file))
avg_line_blogs <- mean(nchar(file))
len_blogs_sample <- length(data_blogs)

# Load the "news" data file
incon <- file(paste(file_path, file_news ,sep=""), "rb")
file <- readLines(incon, encoding="UTF-8", skipNul=TRUE)
data_news <- file[rbinom(length(file), 1, 0.01) == 1]
close(incon)
len_news <- length(file)
size_news <- object.size(file)
max_line_news <- max(nchar(file))
avg_line_news <- mean(nchar(file))
len_news_sample <- length(data_news)

# Load the "twitter" data file
incon <- file(paste(file_path, file_twitter ,sep=""), "r")
file <- readLines(incon, encoding="UTF-8", skipNul=TRUE)
data_twitter <- file[rbinom(length(file), 1, 0.005) == 1]
close(incon)
len_twitter <- length(file)
size_twitter <- object.size(file)
max_line_twitter <- max(nchar(file))
avg_line_twitter <- mean(nchar(file))
len_twitter_sample <- length(data_twitter)

# Now remove the temp file to reselase memory
rm(file)

```

Once the data has been loaded, let's show basic metrics about the data files:

```

# Basic metrics about the files
data_metrics <- data.frame(file_name = c("en_US.blogs.txt", "en_US.news.txt", "en_US.twitter.txt"),
                           size = c(format(size_blogs, units = "auto"),
                                     format(size_news, units = "auto"),
                                     format(size_twitter, units = "auto")),
                           lines = c(format(len_blogs, big.mark=","),
                                     format(len_news, big.mark=","),
                                     format(len_twitter, big.mark=",")),
                           Average_line_length = c(round(avg_line_blogs,0),
                                                    round(avg_line_news,0),
                                                    round(avg_line_twitter,0)),

```

```

max_line_length = c(format(max_line_blogs, big.mark=","),
                     format(max_line_news, big.mark=","),
                     format(max_line_twitter, big.mark=","))
)
# summary table
colnames(data_metrics) <- c('File Name', 'File Size', 'Number of Lines', 'Average Length of Lines', 'Maimun Length of a line')
formattable(data_metrics)

```

File Name

File Size

Number of Lines

Average Length of Lines

Maimun Length of a line

en_US.blogs.txt

255.4 Mb

899,288

230

40,833

en_US.news.txt

257.3 Mb

1,010,242

201

11,384

en_US.twitter.txt

319 Mb

2,360,148

69

140

Data Transformation

On this section I will prepare the data to make it ready for the prediction text model.

To do that:

- Ensure the data is all ASCII.
- The text will be all in lowercase.
- Digits, punctuation, stopwords, white spaces will be removed.
- And the text will be transformred to its stem form.

```

# Make sure all characters are ASCII
data_blogs <- iconv(data_blogs, to="ASCII", sub="")
data_news <- iconv(data_news, to="ASCII", sub="")
data_twitter <- iconv(data_twitter, to="ASCII", sub="")

# Create the corpus to use "tm" packages
my_corpus <- VCorpus(VectorSource(paste(data_blogs, data_news, data_twitter)))

```

```

# Remove original doc to release memory
rm(data_blogs, data_news, data_twitter)

# Transformation using "tm" functions
my_corpus <- tm_map(my_corpus, removeNumbers)
my_corpus <- tm_map(my_corpus, tolower)
my_corpus <- tm_map(my_corpus, removePunctuation)
my_corpus <- tm_map(my_corpus, removeWords, stopwords("english"))
my_corpus <- tm_map(my_corpus, stripWhitespace)
my_corpus <- tm_map(my_corpus, stemDocument)
my_corpus <- tm_map(my_corpus, PlainTextDocument)

# Just to inspect the corpus
inspect(my_corpus[1])

```

```

<<VCorpus>>
Metadata: corpus specific: 0, document level (indexed): 0
Content: documents: 1

```

```

[[1]]
<<PlainTextDocument>>
Metadata: 7
Content: chars: 186

```

```
writeLines(as.character(my_corpus[1]))
```

```

list(list(content = "scene shay apart saturday market park lot logjam driver circl ignor arrow paint pa
list()
list()

```

Data Exploration

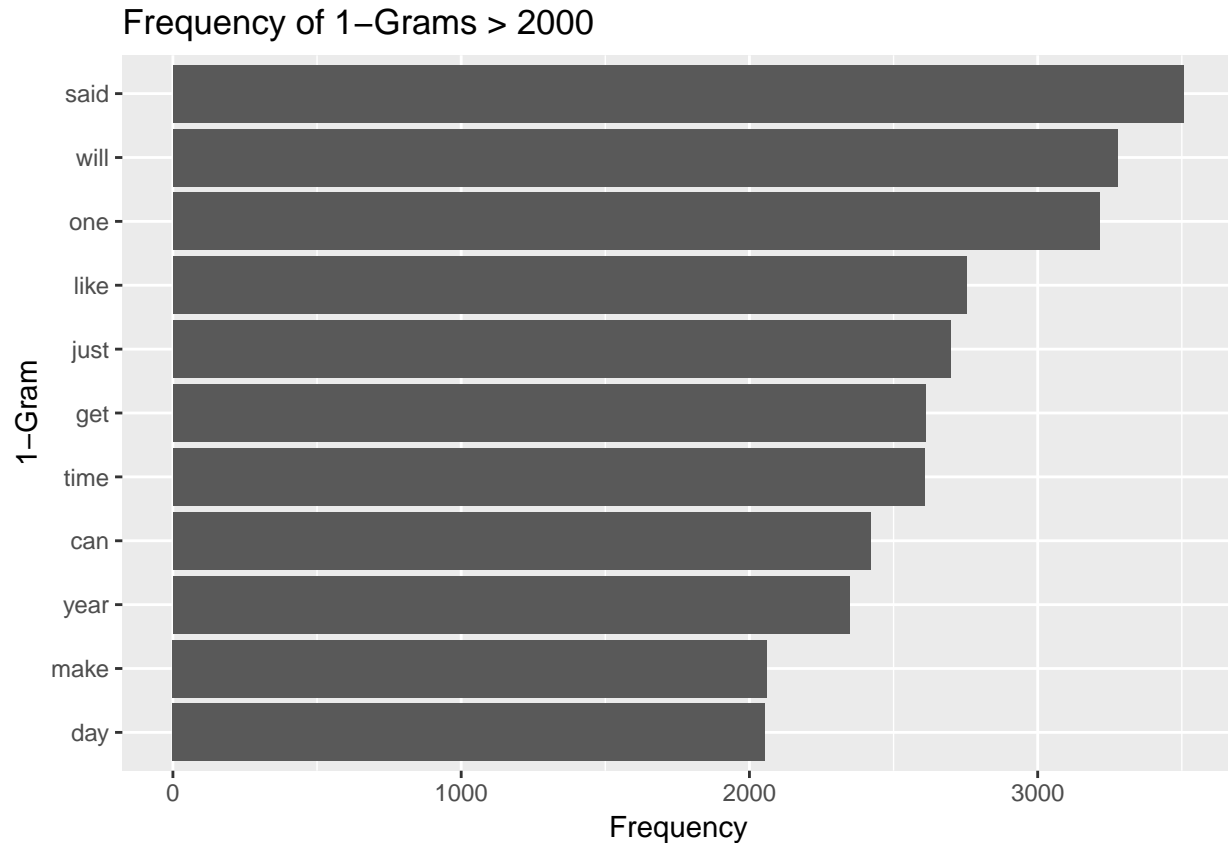
On this section, I will show some data exploration made to the data that was all transformed (previous section) into a **Corpus**.

Let's start showing a **words cloud** with the top 200 most frequent words:

```
wordcloud(my_corpus, max.words = 200, random.order = FALSE, colors=brewer.pal(8,"Dark2"))
```



```
geom_bar(stat = "identity") + coord_flip() +
theme(legend.title=element_blank()) +
xlab("1-Gram") + ylab("Frequency") +
labs(title = "Frequency of 1-Grams > 2000")
```



Bigrams Tokenization:

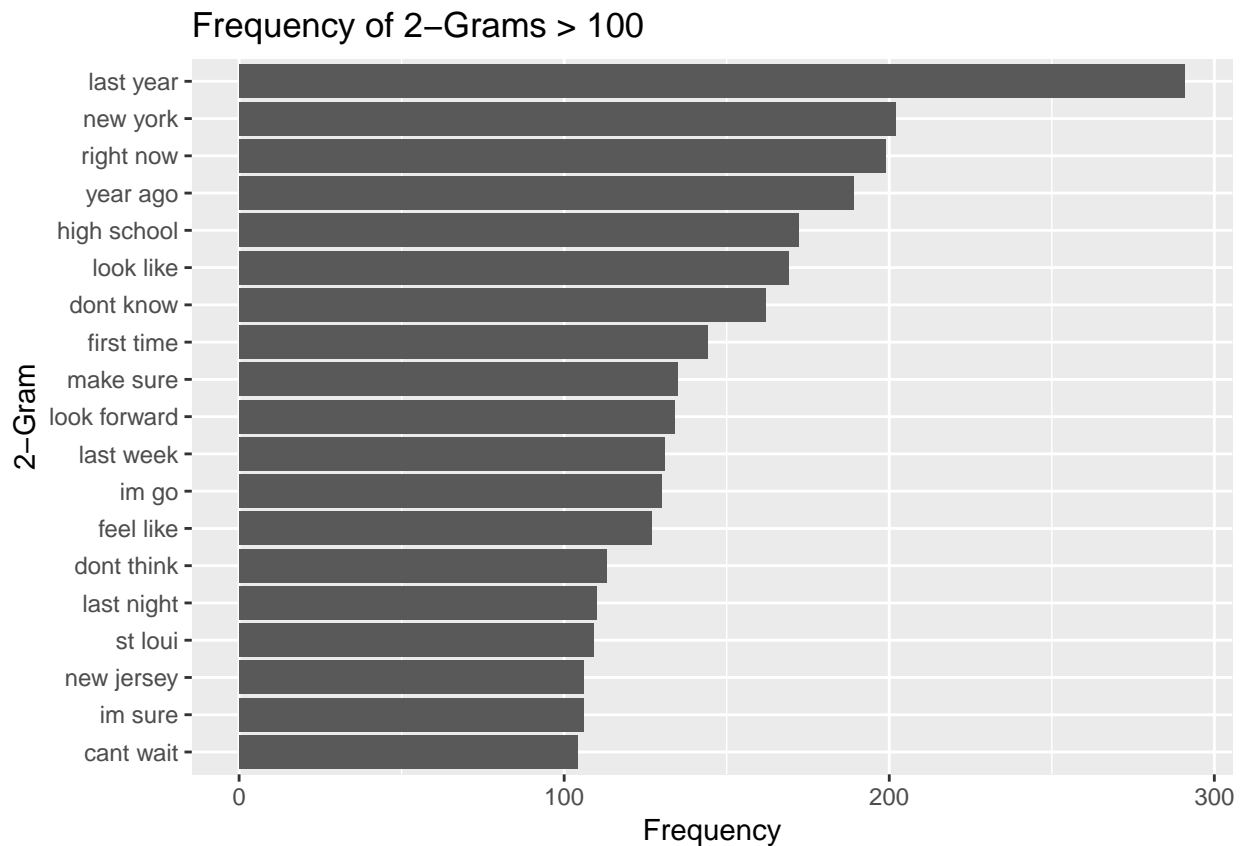
```
two_gramTokenizer <- function(x) NGramTokenizer(x=x, control=Weka_control(min = 2, max = 2))
two_dtm <- DocumentTermMatrix(my_corpus, control = list(tokenize = two_gramTokenizer))
# Because the matrix is too much sparse, let's find something that is not so much
two_dtm_sparse <- removeSparseTerms(two_dtm, sparse=0.9997)

two_dtm_freq <- sort(colSums(as.matrix(two_dtm_sparse)),decreasing = TRUE)
two_dtm_freq_df <- data.frame(word = names(two_dtm_freq), frequency = two_dtm_freq)
head(two_dtm_freq_df, 10)
```

	word	frequency
last year	last year	291
new york	new york	202
right now	right now	199
year ago	year ago	189
high school	high school	172
look like	look like	169
dont know	dont know	162
first time	first time	144
make sure	make sure	135

```
look forward look forward      134
```

```
two_dtm_plot <- subset(two_dtm_freq_df, frequency > 100)
ggplot(two_dtm_plot, aes(x=reorder(word, frequency), y=frequency)) +
  geom_bar(stat = "identity") + coord_flip() +
  theme(legend.title=element_blank()) +
  xlab("2-Gram") + ylab("Frequency") +
  labs(title = "Frequency of 2-Grams > 100")
```



Threegrams Tokenization:

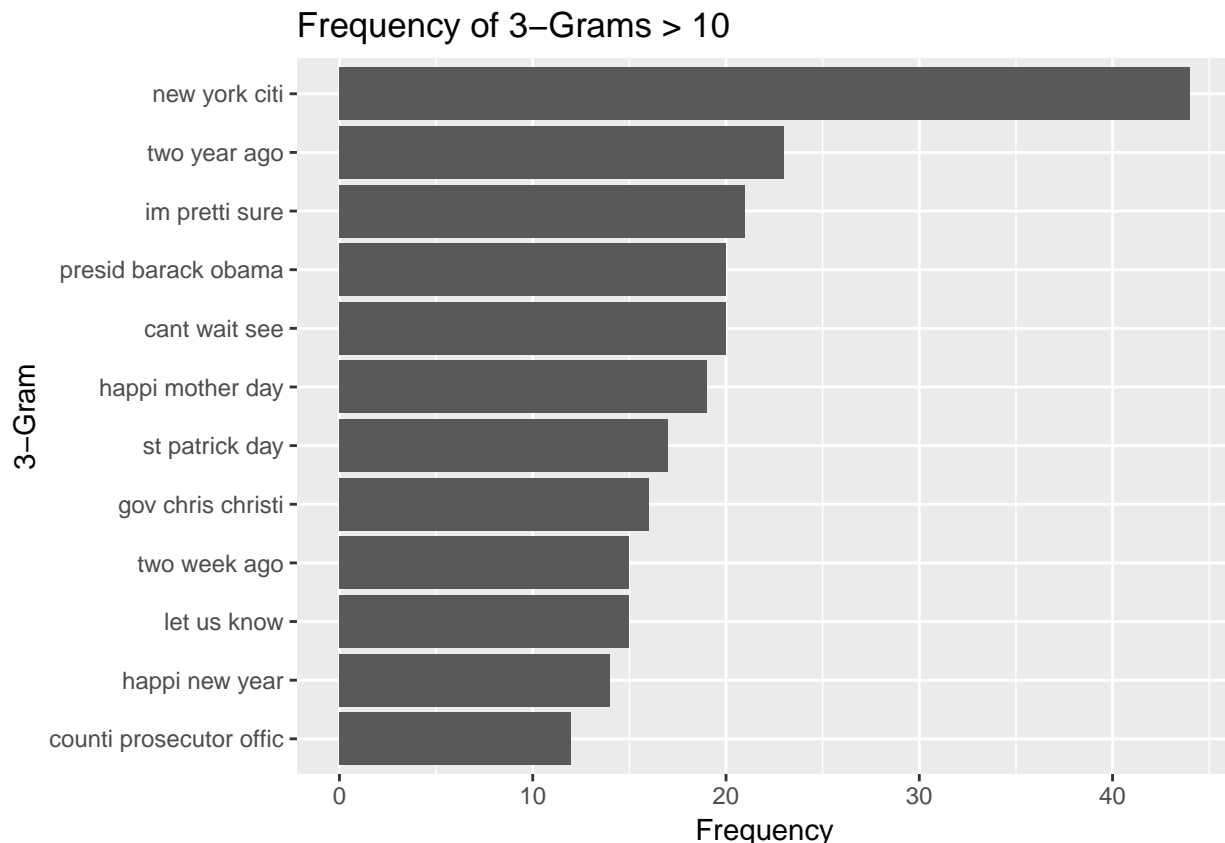
```
three_gramTokenizer <- function(x) NGramTokenizer(x=x, control=Weka_control(min = 3, max = 3))
three_dtm <- DocumentTermMatrix(my_corpus, control = list(tokenize = three_gramTokenizer))
# Because the matrix is too much sparse, let's find something that is not so much
three_dtm_sparse <- removeSparseTerms(three_dtm, sparse=0.9998)

three_dtm_freq <- sort(colSums(as.matrix(three_dtm_sparse)),decreasing = TRUE)
three_dtm_freq_df <- data.frame(word = names(three_dtm_freq), frequency = three_dtm_freq)
head(three_dtm_freq_df, 10)
```

	word	frequency
new york citi	new york citi	44
two year ago	two year ago	23
im pretti sure	im pretti sure	21
cant wait see	cant wait see	20
presid barack obama	presid barack obama	20
happi mother day	happi mother day	19

st patrick day	st patrick day	17
gov chris christi	gov chris christi	16
let us know	let us know	15
two week ago	two week ago	15

```
three_dtm_plot <- subset(three_dtm_freq_df, frequency > 10)
ggplot(three_dtm_plot, aes(x=reorder(word, frequency), y=frequency)) +
  geom_bar(stat = "identity") + coord_flip() +
  theme(legend.title=element_blank()) +
  xlab("3-Gram") + ylab("Frequency") +
  labs(title = "Frequency of 3-Grams > 10")
```



Calcualte of the size of the data-file to be used in the text predition model

This section calculate the size fo the object that hold the ngrams that will be used for the model. Unfortunately there are RAM constraints due the small side of the computer where I'm buiding this project. And based on the guideline of the project, we should assume we will have some RAM constraints on the devide where the model will be running on. So, for those reasons, I'm showing here the size of the ngrams object with some percentages of the corpus:

```
# Basic metrics for the bigrams dataframe
bigrams_metrics <- data.frame(file_name = c("Bigrams 50%", "Bigrams 75%", "Bigrams 90%", "Bigrams 100%"),
  size = c(format(object.size(two_dtm_freq_df[1:round(nrow(two_dtm_freq_df)*0.5),]), units="auto"),
    format(object.size(two_dtm_freq_df[1:round(nrow(two_dtm_freq_df)*0.75),]), units="auto"),
    format(object.size(two_dtm_freq_df[1:round(nrow(two_dtm_freq_df)*0.90),]), units="auto"),
    format(object.size(two_dtm_freq_df[1:round(nrow(two_dtm_freq_df)*1.0),]), units="auto")
  )
```



```

    )

# summary table
colnames(bigrams_metrics) <- c('Bigrams Data', 'File Size')
formattable(bigrams_metrics)

```

Bigrams Data

File Size

Bigrams 50%

1.3 Mb

Bigrams 75%

1.5 Mb

Bigrams 90%

1.7 Mb

Bigrams 100%

1.8 Mb

Strategy to build the Text prediction Model

Now that I have clear understanding of the data, I have transformed it into a clean Corpus, generated the ngrams, and calculated the potential size of the object to be used in the model, follows is the first draft of the strategy on how the text prediction model will work:

- 1- Build the Bigrams with a Corpus Transformed (removed whitespace, stemming, removed stopwords, etc.).
- 2- For quick access build a dictionary with the Bigrams.
- 3- Based on the size constraints, use at least 90% coverage of the Bigrams object.
- 4- When the user enters a word, try to find the word in the dictionary.
- 5- If the word exists, suggest the word that follows based on the Bigrams dataframe.
- 6- If the word does not exist, suggest a word that exists in the Bigrams dataframe selected randomly.

.
 .
 .
 .